

Chapter 7: Synchronization Examples

Outline

- + Explain the bounded-buffer synchronization problem
- + Explain the readers-writers synchronization problem
- + Explain and dining-philosophers synchronization problems
- + Describe the tools used by Linux and Windows to solve synchronization problems.
- + Illustrate how POSIX and Java can be used to solve process synchronization problems

二元號誌 (1/2) - binary semaphore

- ✚ 二元號誌的值只限定為 0 或 1。
 - 利用硬體對二元數值的運算支援，二元號誌的實作要比計數號誌簡單快速得多。
 - 可以利用二元號誌來實作計數號誌。
- ✚ turn, flag, turn+flag
- ✚ test&set, compare&swap, swap
- ✚ mutex lock, semaphore
- ✚ critical region
- ✚ monitor

二元號誌 (2/2)

- 計數號誌可以利用兩個二元號誌以及一個整數實作。

```
void wait(S) {  
    wait(S1);  
    C--;  
    if (C < 0) {  
        signal(S1);  
        wait(S2);  
    }  
    signal(S1);  
}
```

```
void signal(S) {  
    wait( S1);  
    C++;  
    if (C <= 0)  
        signal(S2);  
    else  
        signal(S1);  
}
```

臨界區域 (1/4) - critical region

- ✚ 臨界區域的使用非常方便。
 - 以下宣告一個具有共享變數 v 的臨界區域，在 B 條件式成立下，如果沒有其他行程在此臨界區域中執行，就會執行 S 敘述：

region v when B do S ;

- ✚ 利用臨界區域來實作，程式設計師不用煩惱同步的問題，只要正確地把問題描述在臨界區域內。
- ✚ 有限緩衝區問題可以用臨界區域來簡單地解決同步的問題。

臨界區域 (2/4)

- 生產者與消耗者程式可以分別以臨界區域實作如下。

生產者	消耗者
<pre>region buffer when (count < n) { pool[in] = nextp; in = (in + 1) % n; count++; }</pre>	<pre>region buffer when (count > 0) { nextc = pool[out]; out = (out + 1) % n; count--; }</pre>

臨界區域 (3/4)

- ✦ 臨界區域 **region** v **when B do S** 可利用 **mutex**、**first_delay** 及 **second_delay** 三個號誌實作。
 - **mutex** 號誌是用來確保臨界區的互斥條件成立。
 - 如果行程因為 **B** 為 **FALSE** 而無法進入臨界區，該行程將會在號誌 **first_delay** 等待。
 - 在號誌 **first_delay** 等待的行程重新檢查 **B** 值之前，會離開號誌 **first_delay**，而在號誌 **second_delay** 等待。
 - 分成 **first_delay** 與 **second_delay** 兩段式等待的原因，是為了要避免行程持續忙碌地檢查 **B** 值。
 - 當一個行程離開了臨界區之後，可能因為執行了敘述 **S** 而改變了 **B** 的值，所以需要重新檢查。

臨界區域 (4/4)

```
wait(mutex);
while (!B) {
    first_count++;
    if (second_count > 0)
        signal(second_delay);
    else
        signal(mutex);
    wait(first_delay);
    first_count--;
    second_count++;
    if (first_count > 0)
        signal(first_delay);
    else
        signal(second_delay);
    wait(second_delay);
    second_count--;
}
S;
if (first_count > 0)
    signal(first_delay);
else if (second_count > 0)
    signal(second_delay);
else
    signal(mutex);
```

```
wait(mutex);
while (!B) {
    first_count++;
    if (first_count > 0)
        signal(first_delay);
    else
        signal(mutex);
    —wait(first_delay);
    —first_count--;
    —first_count++;
    —if (first_count > 0)
    —signal(first_delay);
    —else
    —signal(first_delay);
    wait(first_delay);
    first_count--;
}
S;
if (first_count > 0)
    signal(first_delay);
else if (first_count > 0)
    signal(first_delay);
else
    signal(mutex);
```

```
wait(mutex);
while (!B) {
    first_count++;
    if (first_count > 0)
        signal(first_delay);
    else
        signal(mutex);
    wait(first_delay);
    first_count--;
}
S;
if (first_count > 0)
    signal(first_delay);
else
    signal(mutex);
```

```
wait(mutex);
while (!B) {
    first_count++;
    signal(mutex);
    wait(first_delay);
    first_count--;
}
S;
if (first_count > 0)
    signal(first_delay);
else
    signal(mutex);
```


Classical Problems of Synchronization

- + Classical problems used to test newly-proposed synchronization schemes
 - Bounded-Buffer Problem
 - Readers and Writers Problem
 - Dining-Philosophers Problem

Bounded-Buffer Problem

- + N buffers, each can hold one item
- + Semaphore **mutex** initialized to the value 1
- + Semaphore **full** initialized to the value 0
- + Semaphore **empty** initialized to the value N .

Bounded Buffer Problem (Cont.)

producer process

```
do {  
    // produce an item  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
} while (TRUE);
```

consumer process

```
do {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from buffer  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the item  
} while (TRUE);
```

Readers-Writers Problem

- + A data set is shared among a number of concurrent processes
 - **Readers** – only read the data set; they do **not** perform any updates
 - **Writers** – can both read and write
- + Problem – allow multiple readers to read at the same time
 - Only one single writer can access the shared data at the same time
- + Several variations of how readers and writers are considered – all involve some form of priorities
- + Shared Data
 - Data set
 - Semaphore `rw_mutex` initialized to 1
 - Semaphore `mutex` initialized to 1
 - Integer `read_count` initialized to 0

Readers-Writers Problem (Cont.)

- ✚ The structure of a writer process

```
while (true) {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
}
```

Readers-Writers Problem (Cont.)

✚ The structure of a reader process


```
while (true) {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1) /* first reader */  
        wait(rw_mutex);  
    signal(mutex);  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0) /* last reader */  
        signal(rw_mutex);  
    signal(mutex);  
}
```

Readers-Writers Problem Variations

- + The solution in previous slide can result in a situation where a writer process never writes. It is referred to as the “First reader-writer” problem.
- + The “Second reader-writer” problem is a variation the first reader-writer problem that state:
 - Once a writer is ready to write, no “newly arrived reader” is allowed to read.
- + Both the first and second may result in starvation, leading to even more variations
- + Problem is solved on some systems by kernel providing reader-writer locks

Dining-Philosophers Problem

- ✚ N philosophers' sit at a round table with a bowl of rice in the middle.
- ✚ They spend their lives alternating thinking and eating.
- ✚ They do not interact with their neighbors.
- ✚ Occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
- ✚ In the case of 5 philosophers, the shared data
 - ❏ Bowl of rice (data set)
 - ❏ Semaphore chopstick [5] initialized to 1



Dining-Philosophers Problem Algorithm

+ Semaphore Solution

+ The structure of Philosopher i :

```
while (true){  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
  
    /* eat for a while */  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    /* think for a while */  
  
}
```

+ What is the problem with this algorithm?

Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

Solution to Dining Philosophers (Cont.)

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}  
  
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}  
}
```

Solution to Dining Philosophers (Cont.)

- Each philosopher “i” invokes the operations `pickup()` and `putdown()` in the following sequence:

```
DiningPhilosophers.pickup(i);
```

```
/** EAT **/
```

```
DiningPhilosophers.putdown(i);
```

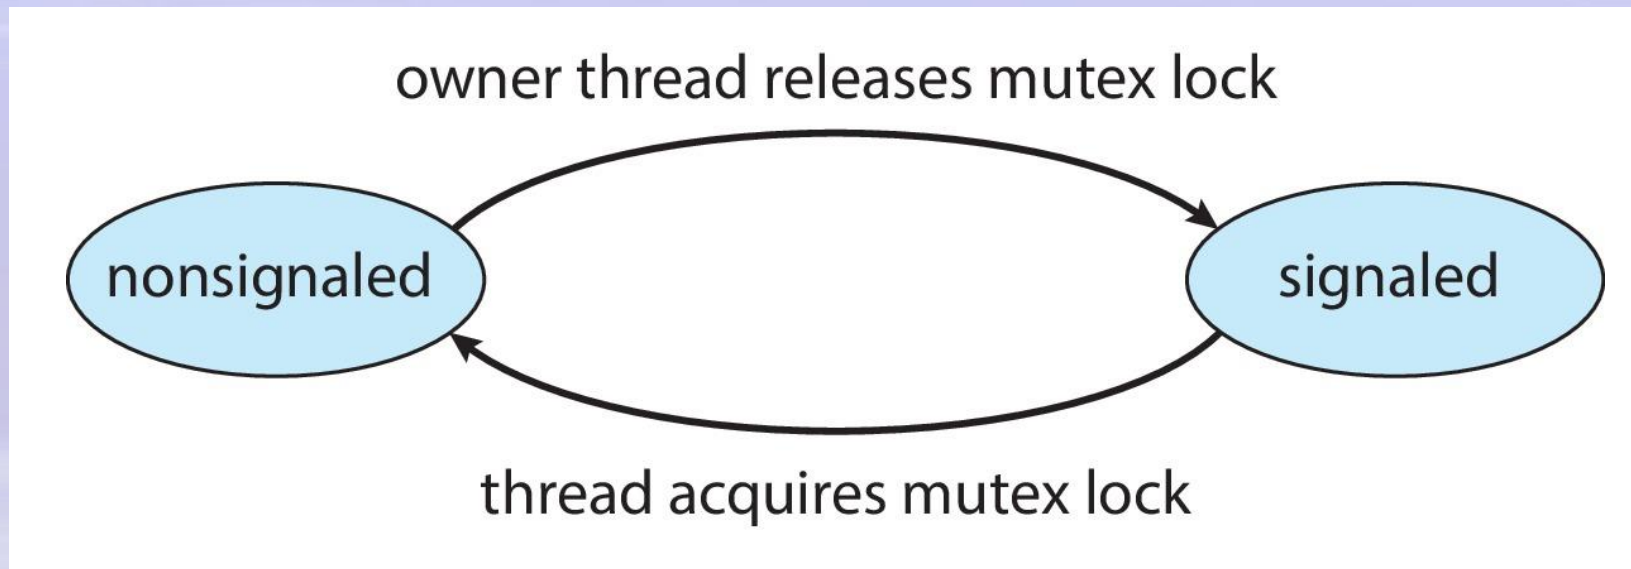
- No deadlock, but starvation is possible

Kernel Synchronization - Windows

- + Uses interrupt masks to protect access to global resources on uniprocessor systems
- + Uses **spinlocks** on multiprocessor systems
 - Spinlocking-thread will never be preempted
- + Also provides **dispatcher objects** user-land which may act mutexes, semaphores, events, and timers
 - **Events**
 - ▣ An event acts much like a condition variable
 - Timers notify one or more thread when time expired
 - Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (thread will block)

Kernel Synchronization - Windows

✚ Mutex dispatcher object



Linux Synchronization

✚ Linux:

- Prior to kernel Version 2.6, disables interrupts to implement short critical sections
- Version 2.6 and later, fully preemptive

✚ Linux provides:

- Semaphores
- Atomic integers
- Spinlocks
- Reader-writer versions of both

✚ On single-CPU system, spinlocks replaced by enabling and disabling kernel preemption

Linux Synchronization

+ Atomic variables

`atomic_t` is the type for atomic integer

+ Consider the variables

```
atomic_t counter;  
int value;
```

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&counter, 5);</code>	<code>counter = 5</code>
<code>atomic_add(10, &counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4, &counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&counter);</code>	<code>value = 12</code>

POSIX Synchronization

- + POSIX API provides
 - mutex locks
 - semaphores
 - condition variable
- + Widely used on UNIX, Linux, and macOS

POSIX Mutex Locks

+ Creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

+ Acquiring and releasing the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

POSIX Semaphores

- ✚ POSIX provides two versions – **named** and **unnamed**.
- ✚ Named semaphores can be used by unrelated processes, unnamed cannot.

POSIX Named Semaphores

- ✚ Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t *sem;

/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```

- ✚ Another process can access the semaphore by referring to its name **SEM**.
- ✚ Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(sem);

/* critical section */

/* release the semaphore */
sem_post(sem);
```


POSIX Unnamed Semaphores

- ✚ Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

- ✚ Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```


POSIX Condition Variables

- Since POSIX is typically used in C/C++ and these languages do not provide a monitor, POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion: Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;  
  
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```

POSIX Condition Variables

- Thread waiting for the condition `a == b` to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
  
pthread_mutex_unlock(&mutex);
```

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

Java Synchronization

- ✚ Java provides rich set of synchronization features:
 - Java monitors
 - Reentrant locks
 - Semaphores
 - Condition variables

Java Monitors

- ✚ Every Java object has associated with it a single lock.
- ✚ If a method is declared as **synchronized**, a calling thread must own the lock for the object.
- ✚ If the lock is owned by another thread, the calling thread must wait for the lock until it is released.
- ✚ Locks are released when the owning thread exits the **synchronized** method.

Bounded Buffer – Java Synchronization

```
public class BoundedBuffer<E>
{
    private static final int BUFFER_SIZE = 5;

    private int count, in, out;
    private E[] buffer;


    public BoundedBuffer() {
        count = 0;
        in = 0;
        out = 0;
        buffer = (E[]) new Object[BUFFER_SIZE];
    }

    /* Producers call this method */
    public synchronized void insert(E item) {
        /* See Slide 37 */
    }

    /* Consumers call this method */
    public synchronized E remove() {
        /* See Slide 37 */
    }
}
```


Java Synchronization

- ✚ A thread that tries to acquire an unavailable lock is placed in the object's **entry set**:



Java Synchronization

- ✚ Similarly, each object also has a **wait set**.
- ✚ When a thread calls **wait()**:
 1. It releases the lock for the object
 2. The state of the thread is set to blocked
 3. The thread is placed in the wait set for the object



Java Synchronization

- ✚ A thread typically calls `wait()` when it is waiting for a condition to become true.
- ✚ How does a thread get notified?
- ✚ When a thread calls `notify()`:
 1. An arbitrary thread T is selected from the wait set
 2. T is moved from the wait set to the entry set
 3. Set the state of T from blocked to runnable.
- ✚ T can now compete for the lock to check if the condition it was waiting for is now true.

Bounded Buffer – Java Synchronization

```
/* Producers call this method */
public synchronized void insert(E item) {
    while (count == BUFFER_SIZE) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    count++;

    notify();
}
```

```
/* Consumers call this method */
public synchronized E remove() {
    E item;

    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;

    notify();

    return item;
}
```

Java Reentrant Locks

- Similar to mutex locks
- The **finally** clause ensures the lock will be released in case an exception occurs in the **try** block.

```
Lock key = new ReentrantLock();  
  
key.lock();  
try {  
    /* critical section */  
}  
finally {  
    key.unlock();  
}
```

Java Semaphores

+ Constructor:

```
Semaphore(int value);
```

+ Usage:

```
Semaphore sem = new Semaphore(1);  
  
try {  
    sem.acquire();  
    /* critical section */  
}  
catch (InterruptedException ie) { }  
finally {  
    sem.release();  
}
```

Java Condition Variables

- + Condition variables are associated with an **ReentrantLock**.
- + Creating a condition variable using **newCondition()** method of **ReentrantLock**:

```
Lock key = new ReentrantLock();  
Condition condVar = key.newCondition();
```

- + A thread waits by calling the **await()** method, and signals by calling the **signal()** method.

Java Condition Variables

- + Example:
- + Five threads numbered 0 .. 4
- + Shared variable `turn` indicating which thread's turn it is.
- + Thread calls `doWork ()` when it wishes to do some work. (But it may only do work if it is their turn.
- + If not their turn, wait
- + If their turn, do some work for awhile
- + When completed, notify the thread whose turn is next.
- + Necessary data structures:

```
Lock lock = new ReentrantLock();  
Condition[] condVars = new Condition[5];  
  
for (int i = 0; i < 5; i++)  
    condVars[i] = lock.newCondition();
```


Java Condition Variables

```
/* threadNumber is the thread that wishes to do some work */
public void doWork(int threadNumber)
{
    lock.lock();

    try {
        /**
         * If it's not my turn, then wait
         * until I'm signaled.
         */
        if (threadNumber != turn)
            condVars[threadNumber].await();

        /**
         * Do some work for awhile ...
         */

        /**
         * Now signal to the next thread.
         */
        turn = (turn + 1) % 5;
        condVars[turn].signal();
    }
    catch (InterruptedException ie) { }
    finally {
        lock.unlock();
    }
}
```


Alternative Approaches

- ✚ Transactional Memory
- ✚ OpenMP
- ✚ Functional Programming Languages

Transactional Memory

- ✚ Consider a function `update()` that must be called atomically. One option is to use mutex locks:

```
void update ()
{
    acquire();

    /* modify shared data */

    release();
}
```

- ✚ A **memory transaction** is a sequence of read-write operations to memory that are performed atomically. A transaction can be completed by adding **`atomic{S}`** which ensure statements in **`S`** are executed atomically:

```
void update ()
{
    atomic {
        /* modify shared data */
    }
}
```

OpenMP

- ✚ OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

- ✚ The code contained within the `#pragma omp critical` directive is treated as a critical section and performed atomically.

Functional Programming Languages

- + Functional programming languages offer a different paradigm than procedural languages in that they do not maintain state.
- + Variables are treated as immutable and cannot change state once they have been assigned a value.
- + There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races.

End of Chapter 7

