

OS Project 2 Report

Design

slave_device.py

The slave device will receive data from master device through socket.

```
1 case slave_IOCTL_MMAP: // add krecv
2     ret = krecv(sockfd_cli, file->private_data, MAP_SIZE, 0);
3     break;
```

```
1 int slave_close(struct inode *inode, struct file *filp)
2 {
3     kfree(filp->private_data); // new
4     return 0;
5 }
6
7 int slave_open(struct inode *inode, struct file *filp)
8 {
9     filp->private_data = kmalloc(MAP_SIZE, GFP_KERNEL); // new
10    return 0;
11 }
```

master_device.py

The master device will allocate some kernel memory to store data and send them to slave device through socket.

```
1 case master_IOCTL_MMAP: // add ksend
2     // send private_data through sockfd_cli
3     ret = ksend(sockfd_cli, file->private_data, ioctl_param, 0);
4     break;
```

```

1  int master_close(struct inode *inode, struct file *filp)
2  {
3      kfree(filp->private_data); // new
4      return 0;
5  }
6
7  int master_open(struct inode *inode, struct file *filp)
8  {
9      filp->private_data = kmalloc(MAP_SIZE, GFP_KERNEL); // new
10     return 0;
11 }

```

Both slave_device.py and master_device.c

We design our `mmap()` function by `remap_pfn_range()` and assign values to `vma`(virtual memory area) attributes.

```

1  static int my_mmap(struct file *filp, struct vm_area_struct *vma)
2  {
3      vma->vm_pgoff = virt_to_phys(filp->private_data)>>PAGE_SHIFT;
4      if(remap_pfn_range(vma, vma->vm_start, vma->vm_pgoff, vma->vm_end-
vma->vm_start, vma->vm_page_prot))
5          return -EAGAIN;
6      vma->vm_flags |= VM_RESERVED;
7      vma->vm_private_data = filp->private_data;
8      vma->vm_ops = &mmap_vm_ops;
9      mmap_open(vma);
10     return 0;
11 }

```

master.c

The master-side program will read the input file with file descriptor `file_fd` by the specified method `mmap` and send the data to the master device.

```

1  case 'm':
2      dst = mmap(NULL, MAP_SIZE, PROT_READ | PROT_WRITE, MAP_SHARED,
dev_fd, 0);
3      for (int j = 0; j * MAP_SIZE < file_size; j++)
4      {
5          // map memory of size MAP_SIZE every loop
6          // if remaining file size is less than MAP_SIZE, then let
mmap_size be the remaining file size
7          mmap_size = (file_size - j * MAP_SIZE > MAP_SIZE)? MAP_SIZE :
file_size - j * MAP_SIZE;
8          src = mmap(NULL, mmap_size, PROT_READ, MAP_SHARED, file_fd, j *
MAP_SIZE);
9          memcpy(dst, src, mmap_size);
10         munmap(src, mmap_size);

```

```

11         while (ioctl(dev_fd, 0x12345678, mmap_size) < 0 && errno ==
EAGAIN);
12     }
13
14     if (ioctl(dev_fd, 0, dst) == -1)
15     {
16         perror("ioctl server error\n");
17         return 1;
18     }
19
20     munmap(dst, MAP_SIZE);
21     break;

```

slave.c

The slave-side program receives the data with size MAP_SIZE each loop from the slave device and writes all of them to the output file by the specified method mmap(). We also use posix_fallocate() to ensure that disk space is allocated for the file referred to by the file descriptor fd for the bytes in the range starting at fileSize and continuing for ret bytes.

```

1  case 'm':
2      src = mmap(NULL, MAP_SIZE, PROT_READ, MAP_SHARED, dev_fd, 0);
3      while (1)
4      {
5          while ((ret = ioctl(dev_fd, 0x12345678)) < 0 && errno == EAGAIN);
6
7          if (ret < 0)
8          {
9              perror("ioctl error\n");
10             return 1;
11         }
12         else if (ret == 0)
13             break;
14         else
15         {
16             offset = (size_t)(file_size / PAGE_SIZE) * PAGE_SIZE;
17
18             posix_fallocate(file_fd, file_size, ret);
19             dst = mmap(NULL, ret, PROT_WRITE, MAP_SHARED, file_fd,
offset);
20             memcpy(dst, src, ret);
21             munmap(dst, ret);
22
23             file_size += ret;
24         }
25     }
26     ftruncate(file_fd, file_size);
27
28     if (ioctl(dev_fd, 0, src) == -1)
29     {
30         perror("ioctl error\n");

```

```

31         return 1;
32     }
33
34     munmap(src, MAP_SIZE);
35     break;

```

Bonus

We modify `kernel.c` which adapts to async IO. We add `FASYNC` to the flags when initializing each socket struct in `ksocket.c` to implement async socket.

```

1  #ifdef ASYNC
2      sk->flags |= FASYNC;
3  #endif

```

Comparison

As the experiments showed below, we can see that transmit file through `mmap` method is faster than I/O method usually because `mmap` only modify the virtual memory mapping but `fcntl` access the disk. However, `fcntl` is slightly faster than `mmap` in the `.iso` case. We guess that different file extensions may cause different ways of allocating/accessing memory so that it is inconsistent with our expected results.

Using async IO, the transmission efficiency has improved in some files, but files like `.mp4` and `.iso` barely improved. We think that async IO may empty the buffer more times than sync IO, which may cause some overhead.

Transmit file	target_file_1 (10 files)	target_file_2	our file (.iso)	our file (.pdf)	our file (.mp4)
sync <code>fcntl</code> IO transmit time	2.531 ms	12.790 ms	19021 ms	25.581 ms	2038 ms
sync <code>mmap</code> IO transmit time	0.586 ms	8.290 ms	19033 ms	11.524 ms	1943 ms
async <code>fcntl</code> IO transmit time	1.812 ms	9.751 ms	19039 ms	13.213 ms	2003 ms

async mmap IO transmit time	0.345 ms	6.344 ms	18961 ms	9.425 ms	1954 ms
--------------------------------------	----------	----------	----------	----------	---------

Member Contribution

B05902010	張頌平	25%
B05902023	李澤諺	25%
B05902040	宋易軒	25%
B05902116	陳昱鈞	25%

Reference

- 1. [mmap\(\)](#)
- 2. [kmalloc\(\)](#) / [kfree\(\)](#)
- 3. [posix_fallocate\(\)](#)